

## Importing the dataset

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.simplefilter("ignore")
```

```
In [2]: ## Displaying the csv file
df=pd.read_csv("heart.csv")
df.head()
```

```
Out[2]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	targ
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	

```
In [3]: df.info()


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   age         1025 non-null   int64  
 1   sex         1025 non-null   int64  
 2   cp          1025 non-null   int64  
 3   trestbps    1025 non-null   int64  
 4   chol        1025 non-null   int64  
 5   fbs         1025 non-null   int64  
 6   restecg     1025 non-null   int64  
 7   thalach     1025 non-null   int64  
 8   exang       1025 non-null   int64  
 9   oldpeak     1025 non-null   float64 
10   slope       1025 non-null   int64  
11   ca          1025 non-null   int64  
12   thal        1025 non-null   int64  
13   target      1025 non-null   int64  
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
In [5]: sns.set_theme(color_codes=True)
```

## Data Visualization


```
In [6]: sns.countplot(data=df,x="cp",hue="target")
```

```
Out[6]: <AxesSubplot:xlabel='cp', ylabel='count'>
```



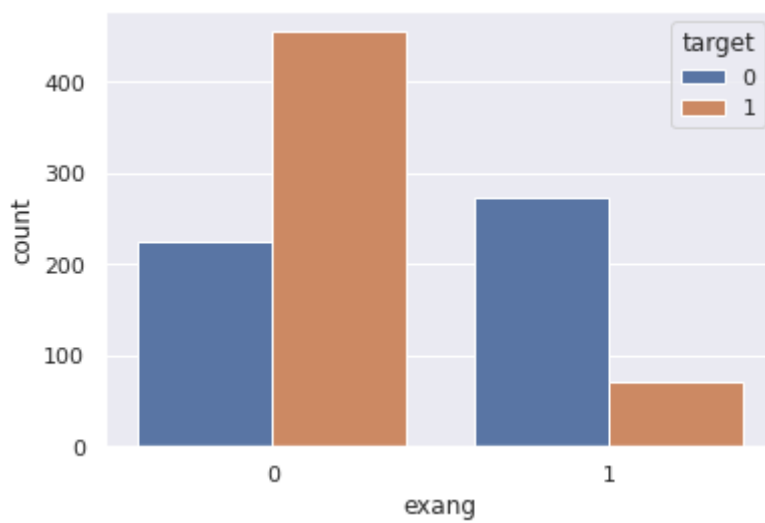
```
In [9]: sns.countplot(data=df,x="fbs",hue="target")
```

```
Out[9]: <AxesSubplot:xlabel='fbs', ylabel='count'>
```




```
In [10]: sns.countplot(data=df,x="exang",hue="target")
```

```
Out[10]: <AxesSubplot:xlabel='exang', ylabel='count'>
```



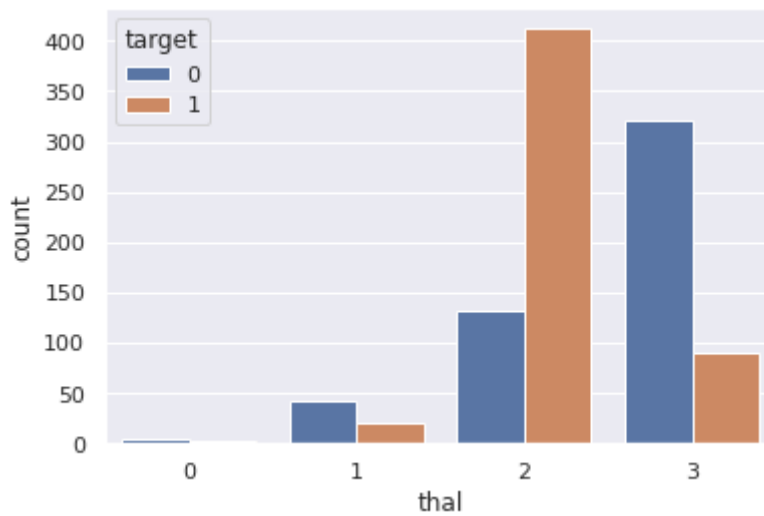
```
In [11]: sns.countplot(data=df,x="ca",hue="target")
```

```
Out[11]: <AxesSubplot:xlabel='ca', ylabel='count'>
```




```
In [12]: sns.countplot(data=df,x="thal",hue="target")
```

```
Out[12]: <AxesSubplot:xlabel='thal', ylabel='count'>
```




```
In [13]: sns.histplot(data=df,x="thal",hue="target")
```

```
Out[13]: <AxesSubplot:xlabel='thal', ylabel='Count'>
```



```
In [14]: sns.histplot(data=df,x="ca",hue="target")
```

```
Out[14]: <AxesSubplot:xlabel='ca', ylabel='Count'>
```



## Data preprocessing

```
In [15]: df.isnull().sum()
```


```
Out[15]: age      0
sex        0
cp         0
trestbps   0
chol       0
fbs        0
restecg    0
thalach    0
exang      0
oldpeak    0
slope      0
ca         0
thal       0
target     0
dtype: int64
```

- We can infer that the dataset does not contain any NAN Values

## Checking if the class label is balanced or not

```
In [18]: sns.countplot(x=df['target'])  
print (df.target.value_counts())
```

```
1    526  
0    499  
Name: target, dtype: int64
```




- We can infer that the class label is not balanced

## Balancing the class label

```
In [24]: from sklearn.utils import resample  
df_majority=df[(df['target']==1)]  
df_minority=df[(df['target']==0)]  
df_minority_upsampled=resample(df_minority,n_samples=526,random_state=42)  
df2=pd.concat([df_majority,df_minority_upsampled])
```

```
In [26]: sns.countplot(x=df2['target'])  
print (df2.target.value_counts())
```


```
1    526  
0    526  
Name: target, dtype: int64
```



## Checking for the impurities


```
In [27]: sns.boxplot(x=df["age"])
```

```
Out[27]: <AxesSubplot:xlabel='age'>
```




```
In [28]: sns.boxplot(x=df["cp"])
```

```
Out[28]: <AxesSubplot:xlabel='cp'>
```



```
In [29]: sns.boxplot(x=df["trestbps"])
```


```
Out[29]: <AxesSubplot:xlabel='trestbps'>
```






```
In [30]: sns.boxplot(x=df["chol"])
```

```
Out[30]: <AxesSubplot:xlabel='chol'>
```




```
In [31]: sns.boxplot(x=df["fbs"])
```

```
Out[31]: <AxesSubplot:xlabel='fbs'>
```



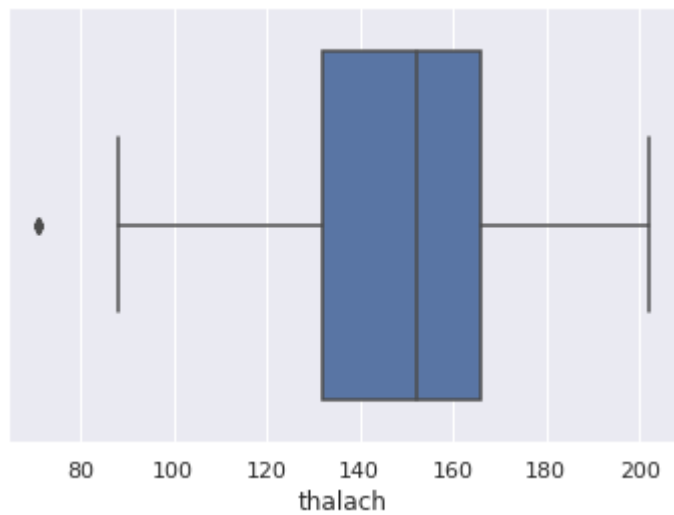
```
In [32]: sns.boxplot(x=df["exang"])
```

```
Out[32]: <AxesSubplot:xlabel='exang'>
```




```
In [33]: sns.boxplot(x=df["thalach"])
```

```
Out[33]: <AxesSubplot:xlabel='thalach'>
```



```
In [34]: sns.boxplot(x=df["restecg"])
```

```
Out[34]: <AxesSubplot:xlabel='restecg'>
```




## Removing the outlier using Z-Score method

```
In [35]: import scipy.stats as stats
z = np.abs(stats.zscore(df2))
data_clean = df2[(z<3).all(axis = 1)]
data_clean.shape
```


```
Out[35]: (989, 14)
```

```
In [36]: sns.heatmap(data_clean.corr(),fmt='.2g')
```

```
Out[36]: <AxesSubplot:>
```



```
In [37]: corr = data_clean[data_clean.columns[1:]].corr()['target'][:-1]
plt.plot(corr)
plt.xticks(rotation=90)
plt.show()
```



## Separating the target variable

```
In [38]: x=data_clean.drop("target",axis=1)
y=data_clean["target"]
```

## Implementing ML Algorithm

```
In [48]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
In [41]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,rand
```

## Logistic Regression

```
In [42]: from sklearn.linear_model import LogisticRegression
```

```
In [43]: model=LogisticRegression()
```

```
In [45]: model.fit(x_train,y_train)
```

```
Out[45]: LogisticRegression()
```

```
In [46]: y_pred_logistic=model.predict(x_test)
```

In [47]: `y_pred_logistic`

Out[47]: `array([[1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,  
1, 1,  
1, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1,  
0, 0,  
1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
1, 1,  
1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1,  
1, 1,  
1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0,  
0, 0,  
1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0,  
0, 1,  
0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0,  
0, 1,  
0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1,  
0, 1,  
0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0,  
1, 1])`

In [49]: `accuracy_logistic=accuracy_score(y_test,y_pred_logistic)*100  
print(f"The accuracy of logistic regression is: {accuracy_logistic}%")`

The accuracy of logistic regression is: 86.86868686868688%

## KNN Algorithm

In [50]: `from sklearn.neighbors import KNeighborsClassifier  
classifier= KNeighborsClassifier(n_neighbors=5, metric='minkowski', p  
classifier.fit(x_train, y_train)  
y_pred_knn= classifier.predict(x_test)  
y_pred_knn`

Out[50]: `array([[0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1,  
0, 1,  
1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1,  
0, 0,  
1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,  
1, 1,  
1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1,  
0, 0,  
1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0,  
0, 0,  
0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0,  
0, 1,  
0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0,  
0, 0,  
0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1,  
0, 0,  
0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0,  
1, 0])`

In [51]: `accuracy_knn=accuracy_score(y_test,y_pred_knn)*100  
print(f"The accuracy of KNN is: {accuracy_knn}%")`

The accuracy of KNN is: 79.7979797979798%

## Random Forest Algorithm

```
In [53]: from sklearn.ensemble import RandomForestClassifier
classifierrandom=RandomForestClassifier(n_estimators=10,criterion="entropy")
classifierrandom.fit(x_train,y_train)
```

```
Out[53]: RandomForestClassifier(criterion='entropy', n_estimators=10)
```

```
In [54]: y_pred_random= classifierrandom.predict(x_test)
y_pred_random
```

```
Out[54]: array([1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1,
1, 1,
1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1,
0, 0,
1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,
1, 1,
1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1,
0, 1,
0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0,
0, 0,
1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0,
0, 1,
1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0,
0, 1,
0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1,
0, 0,
0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0,
1, 0])
```

```
In [55]: accuracy_random=accuracy_score(y_test,y_pred_random)*100
print(f"The accuracy of Random forest algorithm is: {accuracy_random}%")
```

The accuracy of Random forest algorithm is: 98.48484848484848%

## Ada Boost

```
In [56]: from sklearn.ensemble import AdaBoostClassifier
```

```
In [57]: # Create adaboost classifier object
abc = AdaBoostClassifier(n_estimators=50,learning_rate=1)
# Train Adaboost Classifier
modell = abc.fit(x_train, y_train)
```

```
In [58]: y_pred_abc=modell.predict(x_test)
```

```
In [59]: y_pred_abc
```

```
Out[59]: array([[1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1,
 1, 1,
           1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1,
 0, 0,
           1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,
 1, 1,
           1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1,
 1, 1,
           1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0,
 0, 0,
           1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0,
 1, 1,
           0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0,
 0, 1,
           0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1,
 0, 0,
           0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0,
 1, 0])
```

```
In [60]: accuracy_abc=accuracy_score(y_test,y_pred_abc)*100
print(f"The accuracy of Ada Boost algorithm is: {accuracy_abc}%")
```

The accuracy of Ada Boost algorithm is: 89.8989898989899%

## Conclusion

- The accuracy of logistic regression is: 86.86%
- The accuracy of KNN is: 79.79%
- The accuracy of Random forest algorithm is: 98.48%
- The accuracy of Ada Boost algorithm is: 89.89%

```
In [ ]:
```